

Task 3 – One Page Summary

Jurisdiction–Aware Credit Decision Governance Monitor

This project designs a jurisdiction–aware governance monitor for Revolut's AI–assisted credit decisioning. It is built around one product principle: **shared governance schema, not shared sensitive–data pipeline**. The tool uses synthetic credit application data to show how the same credit decisioning problem produces different evidence outputs under US and EU regulatory choices.

Question	Plain–language answer
What the tool does	It monitors AI–assisted credit decision outputs for explanation quality, proxy–risk signals, EU bias–testing metrics, drift and pre–deployment governance evidence, then generates jurisdiction–specific evidence packages.
Why it is needed	A spreadsheet can list rules, but it cannot enforce data–path separation, version rule changes, retain audit logs or route weak evidence to accountable owners. The jurisdiction choice changes data access, computation modules, evidence outputs and escalation owners; it is not just a report–template switch.
How the US and EU differ	US mode focuses on specific adverse–action reason quality, reason–code confidence and proxy–discrimination review. EU mode assumes creditworthiness AI is a high–risk AI use case and checks lifecycle governance evidence, controlled bias testing and post–deployment monitoring.
What data it uses	The prototype uses 1,400 synthetic applications: 700 US processed records and 700 EU audit records. The US processed dataset excludes age, gender_audit, age_band, protected_group_flag and audit_group. The EU audit dataset retains audit–only fields only in a segregated path for controlled bias testing.
What it outputs	US and EU evidence packages use a shared schema: jurisdiction, model version, rule version, risk signal, evidence generated, confidence level, owner, required action, deadline, escalation level and audit reference.
How human oversight works	Human review is risk–tiered. Minor signals create monitoring tickets; moderate issues trigger sampled review; severe governance gaps, sensitive–data leakage or configuration failures trigger model governance review or emergency suspension.
What it deliberately does not do	It does not make final lending decisions, issue legal opinions, certify EU AI Act conformity, automatically repair models or treat SHAP–style attribution as legally sufficient adverse–action reasons. The synthetic figures demonstrate workflow mechanics and are not empirical claims about Revolut customers.